

## Decision Tree

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import plot_tree,export_text
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("data/BikeRental/hour.csv")
df.head()
```

		instant	dte-	sea-	yr	mnth	hr	hol-	week-	work-	weath-	temp	atemp	hum	wind-	ca-	reg-	cnt
			day	son				iday	day	ing-	er-				speed	sua-	stered	
										day					sit			
0	1		2011-01-01 0	1	0	0	6	0	1	0.24	0.287	0.81	0.0	3	13	16		
1	2		2011-01-01 0	1	1	0	6	0	1	0.22	0.272	0.80	0.0	8	32	40		
2	3		2011-01-01 0	1	2	0	6	0	1	0.22	0.272	0.80	0.0	5	27	32		
3	4		2011-01-01 0	1	3	0	6	0	1	0.24	0.287	0.75	0.0	3	10	13		
4	5		2011-01-01 0	1	4	0	6	0	1	0.24	0.287	0.75	0.0	0	1	1		

```
cols_to_drop = ['instant', 'dteday', 'weekday', 'casual', 'registered']
df = df.drop(columns = cols_to_drop)
df.head()
```

		sea-	yr	mnth	hr	holi-	work-	weath-	temp	atemp	hum	wind-	cnt
		son				day	ing-	ersit				speed	
							day						
0	1		0	1	0	0	0	1	0.24	0.2879	0.81	0.0	16
1	1		0	1	1	0	0	1	0.22	0.2727	0.80	0.0	40
2	1		0	1	2	0	0	1	0.22	0.2727	0.80	0.0	32
3	1		0	1	3	0	0	1	0.24	0.2879	0.75	0.0	13
4	1		0	1	4	0	0	1	0.24	0.2879	0.75	0.0	1

```
medCount = df["cnt"].median()
medCount
```

```
142.0
```

We will convert the cnt column, which the number of hirings as categorical as the following:

```
if cnt >= median:
    cnt = 1
else
    cnt = 0
```

```
def cutAtMedian(val):
    return 1 if val >= medCount else 0
```

```
df["cnt"] = df["cnt"].apply(cutAtMedian)
```

```
# save this file as cleaned
df.to_csv("data/BikeRental/cleaned_hour.csv", index=False)
```

```
#loading dataset
def load_dataset(fname):
    data = pd.read_csv(fname, header=0, skip_blank_lines=True)
    dataset = data.values
    X = dataset[:, :-1]
    y = dataset[:, -1]
    cols = data.columns
    return X, y, cols
```

```
#Train Test split
X, y, cols = load_dataset('data/BikeRental/cleaned_hour.csv')
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.33, random_state=1)
print('Train', X_train.shape, y_train.shape)
print('Test', X_test.shape, y_test.shape)
```

```
Train (11643, 11) (11643,)
Test (5736, 11) (5736,)
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```
#Create a decision tree classifier object  
dt = DecisionTreeClassifier(criterion='entropy')
```

```
#Train the decision tree  
dt = dt.fit(X_train, y_train)  
#predict the response for the test data  
yhat=dt.predict(X_test)
```

```
#Evaluate the model  
accuracy=accuracy_score(y_test,yhat)  
print("Accuracy: %.2f" %(accuracy*100))
```

```
Accuracy: 90.66
```


```
plt.figure(figsize=(80,20))  
plot_tree(dt,  
          feature_names=cols[:-1],  
          class_names=['0','1'],  
          max_depth=3  
)
```

```
[Text(2232.0, 978.48, 'hr <= 6.5\nentropy = 1.0\nsamples = 11643\nvalue = [5830, 5813]\nclass = 0'),  
 Text(1116.0, 761.0400000000001, 'hr <= 5.5\nentropy = 0.186\nsamples = 3353\nvalue = [3258, 95]\nclass = 0'),  
 Text(558.0, 543.6, 'hr <= 0.5\nentropy = 0.09\nsamples = 2873\nvalue = [2840, 33]\nclass = 0'),  
 Text(279.0, 326.1600000000001, 'workingday <= 0.5\nentropy = 0.331\nsamples = 476\nvalue = [447, 29]\nclass = 0'),  
 Text(139.5, 108.7200000000003, '\n (...)\n'),  
 Text(418.5, 108.7200000000003, '\n (...)\n'),  
 Text(837.0, 326.1600000000001, 'hr <= 1.5\nentropy = 0.018\nsamples = 2397\nvalue = [2393, 4]\nclass = 0'),  
 Text(697.5, 108.7200000000003, '\n (...)\n'),  
 Text(976.5, 108.7200000000003, '\n (...)\n'),  
 Text(1674.0, 543.6, 'yr <= 0.5\nentropy = 0.555\nsamples = 480\nvalue = [418, 62]\nclass = 0'),  
 Text(1395.0, 326.1600000000001, 'atemp <= 0.538\nentropy = 0.04\nsamples = 232\nvalue = [231, 1]\nclass = 0'),  
 Text(1255.5, 108.7200000000003, '\n (...)\n'),  
 Text(1534.5, 108.7200000000003, '\n (...)\n'),  
 Text(1953.0, 326.1600000000001, 'workingday <= 0.5\nentropy = 0.805\nsamples = 248\nvalue = [187, 61]\nclass = 0'),  
 Text(1813.5, 108.7200000000003, '\n (...)\n'),
```

```

Text(2092.5, 108.72000000000003, '\n (...) \n'),
Text(3348.0, 761.0400000000001, 'temp <= 0.47\nentropy = 0.893\nsamples =
8290\nvalue = [2572, 5718]\nclass = 1'),
Text(2790.0, 543.6, 'yr <= 0.5\nentropy = 0.999\nsamples = 3615\nvalue =
[1868, 1747]\nclass = 0'),
Text(2511.0, 326.1600000000001, 'hr <= 20.5\nentropy = 0.911\nsamples =
1868\nvalue = [1259, 609]\nclass = 0'),
Text(2371.5, 108.72000000000003, '\n (...) \n'),
Text(2650.5, 108.72000000000003, '\n (...) \n'),
Text(3069.0, 326.1600000000001, 'hr <= 21.5\nentropy = 0.933\nsamples =
1747\nvalue = [609, 1138]\nclass = 1'),
Text(2929.5, 108.72000000000003, '\n (...) \n'),
Text(3208.5, 108.72000000000003, '\n (...) \n'),
Text(3906.0, 543.6, 'hr <= 22.5\nentropy = 0.611\nsamples = 4675\nvalue =
[704, 3971]\nclass = 1'),
Text(3627.0, 326.1600000000001, 'hum <= 0.815\nentropy = 0.525\nsamples =
4425\nvalue = [524, 3901]\nclass = 1'),
Text(3487.5, 108.72000000000003, '\n (...) \n'),
Text(3766.5, 108.72000000000003, '\n (...) \n'),
Text(4185.0, 326.1600000000001, 'yr <= 0.5\nentropy = 0.855\nsamples =
250\nvalue = [180, 70]\nclass = 0'),
Text(4045.5, 108.72000000000003, '\n (...) \n'),
Text(4324.5, 108.72000000000003, '\n (...) \n')

```



```
# hyperparameter tuning
dt.tree_.max_depth
```

25

```

acc = []
for max_d in range(1,dt.tree_.max_depth):
    dt=DecisionTreeClassifier(max_depth=max_d,random_state=2)
    dt.fit(X_train,y_train)
    train_accuracy = dt.score(X_train, y_train)
    val_accuracy = dt.score(X_test,y_test)
    acc.append(val_accuracy.round(2))

```

```

print(f'max_depth: {max_d}\t \
      Training Accuracy: {train_accuracy:.2f} \
      Validation accuracy: {val_accuracy:.2f}')

```

max_depth: 1 accuracy: 0.77	Training Accuracy: 0.77	Validation
max_depth: 2 accuracy: 0.81	Training Accuracy: 0.80	Validation
max_depth: 3 accuracy: 0.84	Training Accuracy: 0.84	Validation
max_depth: 4 accuracy: 0.85	Training Accuracy: 0.85	Validation
max_depth: 5 accuracy: 0.88	Training Accuracy: 0.87	Validation
max_depth: 6 accuracy: 0.88	Training Accuracy: 0.88	Validation
max_depth: 7 accuracy: 0.89	Training Accuracy: 0.89	Validation
max_depth: 8 accuracy: 0.89	Training Accuracy: 0.91	Validation
max_depth: 9 accuracy: 0.90	Training Accuracy: 0.93	Validation
max_depth: 10 accuracy: 0.91	Training Accuracy: 0.95	Validation
max_depth: 11 accuracy: 0.91	Training Accuracy: 0.96	Validation
max_depth: 12 accuracy: 0.91	Training Accuracy: 0.97	Validation
max_depth: 13 accuracy: 0.91	Training Accuracy: 0.98	Validation
max_depth: 14 accuracy: 0.91	Training Accuracy: 0.99	Validation
max_depth: 15 accuracy: 0.91	Training Accuracy: 0.99	Validation
max_depth: 16 accuracy: 0.91	Training Accuracy: 0.99	Validation
max_depth: 17 accuracy: 0.90	Training Accuracy: 1.00	Validation
max_depth: 18 accuracy: 0.91	Training Accuracy: 1.00	Validation
max_depth: 19 accuracy: 0.90	Training Accuracy: 1.00	Validation
max_depth: 20 accuracy: 0.90	Training Accuracy: 1.00	Validation
max_depth: 21 accuracy: 0.91	Training Accuracy: 1.00	Validation
max_depth: 22	Training Accuracy: 1.00	Validation

```
accuracy: 0.91          Training Accuracy: 1.00          Validation
max_depth: 23
accuracy: 0.91          Training Accuracy: 1.00          Validation
max_depth: 24
accuracy: 0.91
```

```
max_i = np.argmax(acc)
print(f"The maximum accuracy is at the depth: {max_i+1}
({acc[max_i]*100:.1f}%)")
```

```
The maximum accuracy is at the depth: 10 (91.0%)
```